



# TWEAG

Bayesian data analysis with  
TensorFlow Probability

[www.tweag.io](http://www.tweag.io)

# Your hosts

## Simeon (presentation)



- ▶ background in computational biology
- ▶ Data Scientist at Tweag since 2019
- ▶ lives in Paris

## Dorran (interactive demos & questions in chat)



- ▶ previous positions in geophysics
- ▶ Data Scientist at Tweag since 2019
- ▶ lives in Zurich

Software innovation lab and consultancy based in Paris with employees all around the world and a strong focus on open-source software.

We specialize in

- ▶ software engineering, with a focus on functional programming
- ▶ DevOps, with a focus on reproducible software systems and builds
- ▶ data science

Key industries: finance, biotech, automotive

Need help with your project? Want to work with us?

[www.tweag.io](http://www.tweag.io)

[hello@tweag.io](mailto:hello@tweag.io)

# What you're in for

This tutorial consists of alternating blocks of

- ▶ theory / example slides
  - ▶ practical examples on either external websites or Google Colab notebooks.
- Links at

<https://github.com/tweag/tutorial-dsc-2020>


Requirements:

- ▶ a Google account (for the hands-on demos)
- ▶ elementary knowledge in probability theory and statistics


# Reminder: Probabilities

Probability distributions can be...

discrete:       $\text{Bernoulli}(k; b) = b^k(1 - b)^{1-k}$



continuous:     $\mathcal{N}(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{1}{2\sigma^2}(x - \mu)^2)$



Important concepts:

Conditional probability

$p(A|B)$ : probability that  $A$  is true, given  $B$  is true

Joint probability

$p(A, B)$ : probability that both  $A$  and  $B$  are true

Conditional joint probability

$p(A, B|C)$ : probability that both  $A$  and  $B$  are true, given  $C$  is true

# Bayesian vs frequentist probabilities

Example: fair coin flip with (bias  $b = \frac{1}{2}$ )

## Frequentist probability

$p(\text{"head"}|b = \frac{1}{2})$ :  
 $\frac{\text{\# of heads}}{\text{\# total flips}}$  for  $\infty$  many fair coin flips

## Bayesian probability

$p(\text{"head"}|b = \frac{1}{2})$ :  
measure of *belief* in the statement "flip results in head" given a fair coin

# Prior beliefs

Assume: unknown bias  $b$

## Prior probability


Encodes prior belief in  $b$  *before* flipping the coin

What is known about  $b$ ?

- ▶  $b$  is a probability:  $0 \leq b \leq 1$
- ▶ most coins are fair
- choose prior distribution defined between 0 and 1, with maximum at and symmetric around  $b = \frac{1}{2}$

Example:

$$b \sim \text{Beta}(\alpha = 2, \beta = 2)$$




# Posterior belief

Now: flip coin one time, result: head

## Posterior belief

$p(b|\text{head})$ : updated prior belief after obtaining new data





# Update rule

## Bayes' theorem

$$p(A|B) = \frac{p(B|A) \times p(A)}{p(B)}$$

(easily derived from rules for conditional probabilities)

In data analysis:

$$\underbrace{p(x|D, I)}_{\text{posterior}} = \underbrace{p(D|x, I)}_{\text{likelihood}} \times \underbrace{p(x|I)}_{\text{prior}} / \underbrace{p(D|I)}_{\text{evidence}}$$

$x$ : model parameter (in our case:  $b$ )

$D$ : data (in our case:  $k = 1$ )

$I$ : prior information (often omitted to simplify notation)

# Likelihood


$p(D|x)$ : probability of the data given fixed model parameters  
→ models data-generating process

In our case:

$$p(k|b) = \text{Bernoulli}(k; b) = b^k(1 - b)^{k-1}$$

with

$$k = \begin{cases} 0 : & \text{tail} \\ 1 : & \text{head} \end{cases}$$



# Evidence

$$p(D) = \int dx p(D|x)p(x):$$

normalization constant (long story...)

In our case:

$$\begin{aligned} p(k=1) &= \int_0^1 db L(k=1|b)p(b) \\ &= \int_0^1 db \text{Bernoulli}(k; b) \times \text{Beta}(k; \alpha=2, \beta=2)|_{k=1} \\ &= \int_0^1 db b^k (1-b)^{k-1} \frac{b(1-b)}{\frac{\Gamma(2)\Gamma(2)}{\Gamma(4)}} \bigg|_{k=1} \\ &\vdots \\ &= \frac{1}{2} \end{aligned}$$

# Evidence

$$p(D) = \int dx p(D|x)p(x):$$

normalization constant (long story...)


In our case:

$$\begin{aligned} p(k=1) &= \int_0^1 db L(k=1|b)p(b) \\ &= \int_0^1 db \text{Bernoulli}(k=1|b) \times \text{Beta}(b; \alpha=2, \beta=2)|_{k=1} \\ &= \int_0^1 db b^k (1-b)^{\frac{\Gamma(2)\Gamma(2)}{\Gamma(4)}} \Big|_{k=1} \\ &\vdots \\ &= \frac{1}{2} \end{aligned}$$


# Update rule

In our coin flip example:


prior: 
$$p(b) = \text{Beta}(b; \alpha = 2, \beta = 2)$$
$$\propto b(1 - b)$$



likelihood: 
$$p(D|b) = \text{Bernoulli}(k = 1; b)$$
$$= b$$




posterior: 
$$p(b|D) \propto p(D|b) \times p(b)/p(D)$$
$$= \text{Beta}(b; \alpha = 3, \beta = 2)$$
$$\propto b^2(1 - b)$$



– no slides for interactive stuff –

# Real-world Bayesian data analysis

In real problems: non-standard, difficult posterior distributions



## Probabilistic programming libraries

Allow to

- ▶ programmatically define a statistical model
- ▶ sample from arbitrary posterior distributions
- ▶ run quality checks

Examples:

- ▶ PyMC3
- ▶ Stan
- ▶ TensorFlow Probability
- ▶ ...

# Real-world Bayesian data analysis: recipe

Steps to solve real-world problems:

0. Look at and think about data:

- ▶ What kind of process generated data?
- ▶ What do I want to learn: clusters? Approximating function? class labels?

1. formulate likelihood (statistical model for data-generating process)

2. formulate prior distributions for likelihood parameters


3. set up model in probabilistic programming library

4. sample from posterior distribution


5. perform quality checks



## Step 0: formulating the problem and understanding the data




## Step 0: formulating the problem and understanding the data



Let's do polynomial regression!

## Step 1: formulating the likelihood



Often: measured data = idealized data + noise

idealized data:  $\hat{f}(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$

noise: normally distributed

Likelihood parameters:  $\vec{\beta}, \sigma$

## Step 2: formulating prior distributions

### Incorporate all available information

- ▶ previous experiments
- ▶ ballpark estimation from the data
- ▶ common sense

### Warning

Avoid introducing strong bias!

Noise  $\sigma$ :



- ▶ positive quantity  $\rightarrow p(\sigma)$  with positive support
- ▶ from data: not too big

e.g.,  $\sigma \sim \text{Lognormal}(\mu = 2, \sigma = \frac{1}{2})$

Coefficients  $\beta$ :

- ▶ not too extreme values  $-\infty < \beta_i < \infty$

Something broad and “uninformative”, e.g.  $\mathcal{N}(0, 5)$



## Step 3: code model in probabilistic programming library

– see interactive demo –

## Step 4: sample with appropriate inference algorithm

Most popular and powerful class: Markov chain Monte Carlo (MCMC)

Continuous parameters:

Hamiltonian Monte Carlo (HMC) very efficient

Discrete parameters:

e.g., Metropolis-Hastings

Combination of both:

Gibbs sampling

Another popular option: variational inference (VI)

MCMC introduction blog posts:

[www.tweag.io/blog](http://www.tweag.io/blog)

# Approximation workhorse: Metropolis-Hastings

## Markov chain

Random process with

$$p(x_{i+1}|x_i, x_{i-1}, \dots, x_1) = p(x_{i+1}|x_i)$$

→ a Markov chain has no “memory”

In some conditions: converges to a unique invariant distribution  $\pi(x)$


## Metropolis-Hastings algorithm

Construct Markov chain with invariant distribution  $\pi(x) = p(x)$ :

1. starting at state,  $x_i$ , propose a new state  $x_{i+1}^*$  from  $q(x_{i+1}^*|x_i)$
2. calculate acceptance probability  $p_{\text{acc}}$
3. draw  $u \sim \mathcal{U}(0, 1)$
4. if  $u < p_{\text{acc}}$ :  $x_{i+1} = x_{i+1}^*$ , else  $x_{i+1} = x_i$

# Metropolis (-Hastings)

Initialize with any state  $x_0$




Sequence of states:  
( $x_0$ )



# Metropolis (-Hastings)

Initial state:  $x_0$

1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_0$




# Metropolis (-Hastings)

Initial state:  $x_0$

1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_0$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_1^*)}{p(x_0)} \right)$$



# Metropolis (-Hastings)

Initial state:  $x_0$


1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_0$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_1^*)}{p(x_0)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_1^*$  as the next state  $x_1$ , else copy  $x_0$

Sequence of states:


$(x_0, x_1)$



# Metropolis (-Hastings)

Current state:  $x_1$

1. calculate a proposal state  $x_2^*$  by randomly perturbing  $x_1$




# Metropolis (-Hastings)

Current state:  $x_1$

1. calculate a proposal state  $x_2^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_2^*)}{p(x_1)} \right)$$



# Metropolis (-Hastings)

Current state:  $x_1$


1. calculate a proposal state  $x_2^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_2^*)}{p(x_1)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_2^*$  as the next state  $x_2$ , else copy  $x_1$

Sequence of states:

$(x_0, x_1, x_2)$



# Metropolis (-Hastings)

Current state:  $x_1$


1. calculate a proposal state  $x_2^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_2^*)}{p(x_1)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_2^*$  as the next state  $x_2$ , else copy  $x_1$

Sequence of states:

$$(x_0, x_1, x_2, \dots, x_n)$$



## Step 5: quality checks & debugging

How to check whether sampling was successful?

Effective sample size,  $\hat{R} \rightarrow$  interactive demo

Model debugging strategy:

### Prior predictive check

Use prior samples to generate data points  $y_k$ :

$$p(y_i) = \int d\vec{\beta} d\sigma L(y_k|\vec{\beta}, \sigma) p(\vec{\beta}) p(\sigma)$$

$\rightarrow$  interactive demo

Algorithm:

1. draw samples from prior
2. for each sample  $\beta_i, \sigma_i$ , you get a likelihood  $L(y_k|\beta_i, \sigma_i)$
3. draw samples from  $L(y_k|\beta_i, \sigma_i)$
4. compare if new data points  $y_k$  look  $\approx$  actual data