



# TWEAG

Bayesian data analysis with  
TensorFlow Probability

[www.tweag.io](http://www.tweag.io)

# Your hosts

## Simeon will give the presentation



- ▶ background in computational biology
- ▶ Data Scientist at Tweag since 2019

## Dorran will happily answer questions



- ▶ previous positions in geophysics
- ▶ Data Scientist at Tweag since 2019

## TODO

Tweag I/O is a software innovation lab and consultancy based in Paris with employees all around the world.

We specialize in

- ▶ software engineering, with a focus on functional programming
- ▶ DevOps, with a focus on reproducible software systems and builds
- ▶ data science

# What you're in for

This tutorial consists of alternating blocks of

- ▶ theory / example slides
- ▶ practical examples on either external websites or Google Colab notebooks.  
Links are provided at <https://github.com/tweag/tutorial-dsc-2020/>

Requirements:


- ▶ a Google account (for the practical exercises)
- ▶ elementary knowledge in probability theory and statistics

# Reminder: Probabilities

Probability distributions can be...

discrete:  $\text{Ber}(k; b) = b^k(1 - b)^{1-k}$

continuous:  $\mathcal{N}(\mathbf{x}; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{1}{2\sigma^2}(\mathbf{x} - \mu)^2)$



Important concepts:

Conditional probability

$p(A|B)$ : probability that  $A$  is true, given  $B$  is true

Joint probability

$p(A, B)$ : probability that both  $A$  and  $B$  are true

Conditional joint probability

$p(A, B|C)$ : probability that both  $A$  and  $B$  are true, given  $C$  is true

$p(x)$

$p(x)$

# Bayesian vs frequentist probabilities

Example: fair coin flip with  $p(\text{head}) = p(\text{tail}) = \frac{1}{2}$

## Frequentist probability

$$p(\text{"flip results in head"}|b = \frac{1}{2}) = \frac{1}{2}:$$

$\frac{\text{\# of heads}}{\text{\# total flips}}$  for  $\infty$  many fair coin flips

## Bayesian probability

$$p(\text{"flip results in head"}|b = \frac{1}{2}) = \frac{1}{2}:$$

measure of *belief* in the statement "flip results in head" given single fair coin flip

**TODO:** this slide looks complicated and is aesthetically offputting

# Prior beliefs

Assume: unknown bias  $b$

## Prior probability

Encodes prior belief in  $b$  *before* flipping the coin

What is known about  $b$ ?

- ▶  $b$  is a probability:  $0 \leq b \leq 1$
- ▶ most coins are fair
- choose prior distribution defined between 0 and 1, with maximum at and symmetric around  $b = \frac{1}{2}$

Example:

$$b \sim \text{Beta}(\alpha = 2, \beta = 2)$$

with


$$\text{Beta}(x; \alpha, \beta) \propto x^{\alpha-1} (1-x)^{\beta-1}$$

# Posterior belief

Now: flip coin one time, result: head

## Posterior belief

$p(b|\text{head})$ : updated prior belief after obtaining new data





# Update rule

## Bayes' theorem

$$p(A|B) = \frac{p(B|A) \times p(A)}{p(B)}$$

(easily derived from rules for conditional probabilities)

In data analysis:

$$\underbrace{p(x|D, I)}_{\text{posterior}} = \underbrace{p(D|x, I)}_{\text{likelihood}} \times \underbrace{p(x|I)}_{\text{prior}} / \underbrace{p(D|I)}_{\text{evidence}}$$

$x$ : model parameter

$D$ : data

$I$ : prior information (often not made explicit)

# Likelihood


$p(D|x)$ : probability of the data given fixed model parameters  
→ models data-generating process

In our case:

$$p(k|b) = b^k(1-b)^{k-1} = \text{Ber}(k; b)$$

with

$$k = \begin{cases} 0 : & \text{tail} \\ 1 : & \text{head} \end{cases}$$



# Evidence

$$p(D) = \int dx p(D|x)p(x):$$

normalization constant (long story...)

In our case:

$$\begin{aligned} p(k=1) &= \int_0^1 db L(k=1|b)p(b) \\ &= \int_0^1 db \text{Ber}(k;b) \times \text{Beta}(k; \alpha=2, \beta=2)|_{k=1} \\ &= \int_0^1 db b^k(1-b)^{k-1}b(1-b)|_{k=1} \\ &\vdots \\ &= \frac{1}{12} \end{aligned}$$

# Evidence

$$p(D) = \int dx p(D|x)p(x):$$

normalization constant (long story...)


In our case:

$$\begin{aligned} p(k=1) &= \int_0^1 db L(k=1|b)p(b) \\ &= \int_0^1 db \text{Ber}(k;b) \text{Beta}(k; \alpha=2, \beta=2)|_{k=1} \\ &= \int_0^1 db b^1(1-b)^{2-1} b(1-b)|_{k=1} \\ &\vdots \\ &= \frac{1}{12} \end{aligned}$$

# Update rule

In our coin flip example:

prior: 
$$p(b) = \text{Beta}(b; \alpha = 2, \beta = 2)$$
$$\propto b(1 - b)$$



likelihood: 
$$p(D|b) = \text{Ber}(k = 1; b)$$
$$= b$$




posterior: 
$$p(b|D) \propto p(D|b) \times p(b)/p(D)$$
$$= \text{Beta}(b; \alpha = 3, \beta = 2)$$
$$\propto b^2(1 - b)$$



– no slides for interactive stuff –

# Real-world-ish application: regression

**Step 1:** postulate likelihood  $p(D|x)$  or: think about and look at the data



Standard distributions generate certain types of data:

**Poisson distribution:** counts per interval

**Exponential distribution:** interarrival times

...

Often: measured data = idealized data + noise


We guess:

**idealized data:**  $\hat{f}(x; \vec{\beta}) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$

**noise:** normally distributed

# Real-world-ish application: regression

## Step 2: find sensible prior distributions



Incorporate all available information:

- ▶ noise  $\sigma$  is positive quantity  $\rightarrow$  choose  $p(\sigma)$  with positive support
- ▶ eyeballing the data
- ▶ common sense
- ▶ ...

**Warning:** avoid introducing strong bias!

Principled methods are available to find good prior distributions.

In our case:

$\sigma$ : for example,  $\sigma \sim \text{Lognormal}(\mu = 2, \sigma = \frac{1}{2})$

$\vec{\beta}$ : broad, rather uninformative normal distributions, e.g.,  $\mathcal{N}(0, 5)$



# Real-world-ish application: regression

**Step 3:** specify model in probabilistic programming library

- ▶ programmatic formulation of statistical model
- ▶ powerful inference algorithms
- ▶ “debugging” functionality

Popular and easy-to-use choices are PyMC3, Stan, or TensorFlow Probability.

**Step 4:** perform inference with appropriate algorithm

Most popular and powerful class: Markov chain Monte Carlo

**Continuous parameters:** Hamiltonian Monte Carlo (HMC) very efficient

**Discrete parameters:** e.g., Metropolis-Hastings

**Combination of both:** Gibbs sampling

Another popular option: variational inference (VI)




# Real-world issues

In reality, probability distributions often

- ▶ of non-standard form
- ▶ are multidimensional
- ▶ have highly correlated random variables
- ▶ are known only up to a normalization constant

Consequences:

- ▶ analytical evaluation of expectation values is impossible
- ▶ naïve sampling approaches are inefficient (curse of dimensionality)



# Approximation workhorse: Metropolis-Hastings

## Markov chain

Random process with

$$p(x_{i+1}|x_i, x_{i-1}, \dots, x_1) = p(x_{i+1}|x_i)$$

→ a Markov chain has no “memory”

In some conditions: converges to a unique invariant distribution  $\pi(x)$


## Metropolis-Hastings algorithm

Construct Markov chain with invariant distribution  $\pi(x) = p(x)$ :

1. starting at state,  $x_i$ , propose a new state  $x_i^*$  from  $q(x_i^*|x_i)$
2. calculate acceptance probability  $p_{\text{acc}}$
3. draw  $u \sim \mathcal{U}(0, 1)$
4. if  $u < p_{\text{acc}}$ :  $x_{i+1} = x_i^*$ , else  $x_{i+1} = x_i$

# Metropolis(-Hastings)

Initialize with any state  $x_0$




Sequence of states:  
( $x_0$ )

# Metropolis(-Hastings)

Initial state:  $x_0$

1. calculate a proposal state  $x_0^*$  by randomly perturbing  $x_0$




# Metropolis(-Hastings)

Initial state:  $x_0$

1. calculate a proposal state  $x_0^*$  by randomly perturbing  $x_0$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_0^*)}{p(x_0)} \right)$$



# Metropolis(-Hastings)

Initial state:  $x_0$


1. calculate a proposal state  $x_0^*$  by randomly perturbing  $x_0$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_0^*)}{p(x_0)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_0^*$  as the next state  $x_1$ , else copy  $x_0$

Sequence of states:


$(x_0, x_1)$



# Metropolis(-Hastings)

Current state:  $x_1$

1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_1$






# Metropolis(-Hastings)

Current state:  $x_1$

1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_1^*)}{p(x_1)} \right)$$



# Metropolis(-Hastings)

Current state:  $x_1$


1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_1^*)}{p(x_1)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_1^*$  as the next state  $x_2$ , else copy  $x_1$

Sequence of states:

$(x_0, x_1, x_2)$



# Metropolis(-Hastings)

Current state:  $x_1$

1. calculate a proposal state  $x_1^*$  by randomly perturbing  $x_1$
2. calculate acceptance probability

$$p_{\text{acc}} = \min \left( 1, \frac{p(x_1^*)}{p(x_1)} \right)$$

3. with probability  $p_{\text{acc}}$ , accept proposal state  $x_1^*$  as the next state  $x_2$ , else copy  $x_1$

Sequence of states:

$(x_0, x_1, x_2, \dots, x_n)$

